



## Preparation Exam: Certified Developer - Associate for AWS (June 2018) Exam Session - Failed

[Repeat session](#)

### Overall Results

Share on: [Twitter](#) [Facebook](#) [LinkedIn](#) [Google+](#)

#### Score

55/100

#### Answers

Correct	55%
Incorrect	45%
Skipped	0%




#### Time spent

1h 45m

### Progress

Overall



	Score	Rank
Previous Score	192	1084th
This Exam	224 +32	905th

Assessed skills

Your score has been updated for the following skills

Security for AWS	530
<a href="#">Improve your score</a> <a href="#">Browse Training</a>	+236
Development for AWS NEW	433
<a href="#">Improve your score</a> <a href="#">Browse Training</a>	
Deployment for AWS NEW	348
<a href="#">Improve your score</a> <a href="#">Browse Training</a>	
Serverless for AWS	490
<a href="#">Improve your score</a> <a href="#">Browse Training</a>	+277

This assessment also provided us feedback on

AWS	288
Security	▲ 291
Analytics	▲ 130
Networking	201

[Show More](#)

Areas

☐ Deployment

Your score: 64%

Learning Path

Required score: 35% <input type="checkbox"/>
<div><input type="checkbox"/> <b>Development</b></div> <div>Your score: 50%</div> <div>Required score: 35% <input type="checkbox"/></div>
<div><input type="checkbox"/> <b>Refactoring</b></div> <div>Your score: 86%</div> <div>Required score: 35% <input type="checkbox"/></div>
<div><input type="checkbox"/> <b>Monitoring and Troubleshooting</b></div> <div>Your score: 13%</div> <div>Required score: 35% <input type="checkbox"/></div>
<div><input type="checkbox"/> <b>Security</b></div> <div>Your score: 63%</div> <div>Required score: 35% <input type="checkbox"/></div>

Review answers

<u>Wrong</u>	29
<u>Correct</u>	36

#1

KMS

Which statement is **false** about Key Admin Policy with regard to Key Management Service?

✓ Key Admin Policy is defined and maintained in KMS, and is isolated from IAM policies.

B Key Admin Policy can manage access based on users and roles, but cannot manage access based on groups.

C Key Admin Policy is defined and maintained in IAM Policies; any group with this policy attached will be able to administrate keys in KMS.

D Key Admin Policy uses the principal object to identify users or roles that policy applies to.

#### Explanation

Key Admin Policy is created and managed in KMS, not in IAM, and uses the Principal object to identify users or roles that the policy applies to. Key Admin Policy can not be applied to groups.

<https://cloudacademy.com/amazon-kms-key-admin-policy/>

## #3

Kinesis Lambda

Your company is using Amazon Kinesis and would like some custom functionality when certain events occur. As a result, they are thinking of integrating AWS Lambda with Amazon Kinesis and would like you to give a presentation on how integration works. What might be good topics to discuss? (Choose 2 answers)

✓ That the integration will be a stream-based model where AWS Kinesis polls the Amazon Lambda stream.

B That the integration will be a stream-based model where AWS Lambda polls the Amazon Kinesis stream.

✓ That AWS Lambda would need to be synchronously invoked.

D That AWS Lambda would need to be asynchronously invoked.

#### Explanation

Kinesis and Lambda integration functions using a stream-based model where AWS Lambda polls the stream and when it detects new records invokes your Lambda function by passing the new records as a parameter. In terms of invocation type, AWS Lambda invokes a Lambda function using the RequestResponse invocation type (synchronous invocation) by polling the Kinesis Stream.

<http://docs.aws.amazon.com/lam...>

#6

You've deployed an application in a custom AMI image into the Amazon cloud. It is deployed in a separate VPC. You would like to take advantage of being able to failover to another instance without having to reconfigure the application. Which of these solutions could be utilized? (Choose 2 answers)

A Use an additional elastic network interface for failover to another instance.

✓ Utilize Cloud Watch health checks for failover.

✓ Use load balancing to balance traffic to additional instances.

D Add a secondary private IP address to the primary network interface that could then be used to move to a failover instance.

#### Explanation

The ENI can only be attached to an instance hosted in a VPC. When you move a network interface from one instance to another, network traffic is redirected to the new instance. Some network and security appliances, such as load balancers, network address translation (NAT) servers, and proxy servers prefer to be configured with multiple network interfaces. You can create and attach secondary network interfaces to instances in a VPC that are running these types of applications and configure the additional interfaces with their own public and private IP addresses, security groups, and source/destination checking.

<http://docs.aws.amazon.com/AWS...>

#9

AWS CodeCommit

You have migrated to AWS CodeCommit, but the entire push seems to be failing due to intermittent network issues. What could be a possible solution to this?

A Delete the push history from your repository and retry.

✓ Use SSH to push your repository over the network.

C Repush the entire repository.

D Push your repository in increments or chunks.

#### Explanation

When migrating to AWS CodeCommit, it is recommended to push your repository in increments or chunks to reduce the chances of intermittent network issues or degraded network performance causing the entire push to fail.

<http://docs.aws.amazon.com/cod...>

## #11

CloudFormation

You are replicating your company's AWS infrastructure in another region, creating a disaster recovery environment. To achieve this, you modify your CloudFormation template to use the new DR region. However, when you run the deployment, it fails. What are some possible reasons for this? (Choose 3 answers)

✓ The disaster recovery region has a different number of availability zones.

✓ You need to copy your CloudFormation template to the disaster recovery region.

C Resource identifiers such will be different from region to region.

✓ EC2 key pairs will differ from region to region.

#### Explanation

Resource ID's such as subnet and AMI id's, as well as key pairs, will differ from region to region. You can use the GetAZ function to list all available zones within a region.

<http://docs.aws.amazon.com/AWS...>

## #17

EC2 VPC

You want to connect to an EC2 instance inside your VPC. Unfortunately, you are getting the error "Network error: Connection timed out." What checks might you perform to troubleshoot the issue? (Choose 3 answers)

✓ Check the permissions on your IAM policy.

✓ Check your security group rules.

✓ Check the network ACL for your subnet.

D Check the route table for your subnet.

**Explanation**

If you're having a network connection and you're using a VPC, it's a good idea to make sure that your VPC subnet is configured to route all traffic destined outside the VPC to the Internet gateway for the VPC (check the route table). It's also a good idea to make sure that the network ACLs is allowing inbound and outbound traffic from your local IP address on the proper port. Making sure you have a security group rule that allows inbound traffic from your public IPv4 address on the proper port is also important.

<http://docs.aws.amazon.com/AWS...>

## #18

CloudWatch EC2

A new client has an application replicas hosted Amazon EC2 instances deployed in different regions. They want to be able to use CloudWatch to monitor and analyze all of these EC2 instances efficiently. What recommendations can you make? (Choose 3 answers)

A Pull and output the appropriate data from each region using the CloudWatch command line interface.

✓ Note that Amazon CloudWatch cannot aggregate data across regions but that regional monitoring is possible.

✓ Select the EC2 namespace in the CloudWatch console and then select Across All Instances to graph each region.

✓ Set a CloudWatch alarm indicating when the data is ready for aggregation.

### Explanation

You can select metrics and create graphs of the data using the CloudWatch console, and view your data at different granularities. For example, you can choose a detailed view (for example 1 minute), which can be useful when troubleshooting. You can choose a less detailed view (for example, 1 hour), which can be useful when viewing a broader time range (for example, 3 days) so that you can see trends over time.

Note that you can aggregate only statistics across regional EC2 instances that have detailed monitoring enabled. Instances that use basic monitoring are not included. Therefore, you must enable detailed monitoring (at an additional charge), which provides data in 1-minute periods.

<http://docs.aws.amazon.com/Ama...>

## #24

Lambda API Gateway

You are coming on as a developer for an Android mobile application your company runs. You notice that all Lambda invocations are made over HTTP using the Amazon API Gateway instead of through the AWS Android SDK. What desirable objectives does this design achieve? (Choose 3 answers)

A It provides the ability to throttle individual users or requests.

✓ It enables the default connection and settings for Lambda integration with Amazon Cognito.

✓ It provides a caching layer to store responses from Lambda functions.

✓ It adds automatic protection against Distributed Denial of Service attacks.

### Explanation



You can invoke your Lambda function over HTTP using Amazon API Gateway (instead of using any of the AWS SDKs). Amazon API Gateway adds an additional layer between your mobile users and your app logic that enable the following:

- Ability to throttle individual users or requests
- Protect against Distributed Denial of Service attacks
- Provide a caching layer to cache response from your Lambda function.

<http://docs.aws.amazon.com/lam...>

## #25

IAM

You are working as a developer on your client's cloud deployment trying to federate with Active Directory. Specifically, you are testing the API call `AssumeRoleWithSAML` to get temporary security credentials for the user. What security credentials should be returned by this call? (Choose 3 answers)

✓ Encryption key

B Security token

✓ Security access key

✓ Access key ID

### Explanation

`AssumeRoleWithSAML` returns a set of temporary security credentials for users who have been authenticated via a SAML authentication response. The temporary security credentials returned by this operation consist of an access key ID, a secret access key, and a security token. Applications can use these temporary security credentials to sign calls to AWS services.

<http://docs.aws.amazon.com/STS...>

## #29

API Gateway

You are developing an API in Amazon API Gateway that several mobile applications will use to interface with a back end service in AWS being written by another developer. You can use a(n)\_\_\_ integration for your API methods to develop and test your client applications before the other developer has completed work on the back end.

✓ Lambda function

B AWS service proxy

C Mock

D HTTP proxy

#### Explanation

Amazon API Gateway supports mock integrations for API methods. This feature enables API developers to generate API responses from API Gateway directly, without the need for an integration back end. As an API developer, you can use this feature to unblock other dependent teams needing to work with an API before the project development is complete. You can also leverage this feature to provision a landing page of your API, which can provide an overview of and navigation to your API.

<http://docs.aws.amazon.com/api...>

## #30

DynamoDB

Consider a DynamoDB table used to track users registered to a blog site. Each registered user results in an item with the following attributes: UserID, Name, Email, and RegistrationDate. Given that you'll need to query for newly registered users, what is the most appropriate key structure to use for this table?

A RegistrationDate as the Partition Key

✓ UserID as the Partition Key

C UserID as the Partition Key; RegistrationDate as the Sort Key

D RegistrationDate as the Partition Key; UserID as the Sort Key

#### Explanation

A composite key using UserID as the partition key and RegistrationDate as the sort key would be most appropriate to return a sorted list of registered users by date.

<http://docs.aws.amazon.com/ama...>

## #35

Direct Connect

How does AWS Direct Connect differ from a VPN Connection?

- A AWS Direct Connect can be configured in minutes.
- ✓ AWS Direct Connect utilizes IPSec to establish encrypted network connectivity between your intranet and Amazon VPC over the Internet.
- C AWS Direct Connect can tolerate the inherent variability in Internet-based connectivity.
- D AWS Direct Connect uses dedicated, private network connections between your intranet and Amazon VPC.

**Explanation**

A VPC VPN Connection utilizes IPSec to establish encrypted network connectivity between your intranet and Amazon VPC over the Internet. VPN Connections can be configured in minutes and are a good solution if you have an immediate need, have low to modest bandwidth requirements, and can tolerate the inherent variability in Internet-based connectivity.

**AWS Direct Connect does not involve the Internet;** instead, it uses dedicated, private network connections between your intranet and Amazon VPC.

<http://aws.amazon.com/directco...>

## #37

CloudFormation

You have been assigned to a client to use OpsWorks for deployments in their AWS cloud environment. The client is doing web development and needs environments for dev, test, and production. But the main requirement you have been given is to control cost wherever possible. How can you configure and deploy these environments using OpsWorks while controlling costs?

- A Get pre-configured Chef recipes at reduced cost.

B Auto scaling has a charge in a VPC but is not charged when configured using CloudFormation.

✓ Using CloudFormation reduces the costs of EC2 instances.

D Use different CloudFormation templates for dev, test, and production to control the size and cost of servers deployed.

#### Explanation

To provision and configure your stack resources, you must understand AWS CloudFormation templates, which are formatted text files in JSON or YAML. These templates describe the resources that you want to provision in your AWS CloudFormation stacks. You can use AWS CloudFormation Designer or any text editor to create and save templates. You can use different templates for dev, test, and production environments which can save cost in dev and test. If you are doing performance testing you can use a template to quickly set up a full size environment and take down when testing is completed.

<http://docs.aws.amazon.com/AWS...>

## #40

S3

A storage admin wants to encrypt all the objects stored in S3 using server-side encryption. The user does not want to use the AES-256 encryption key provided by S3. How can the user achieve this?

✓ The admin should use CLI or API to upload the encryption key to the S3 bucket. When making a call to the S3 API mention the encryption key URL in each request

B The admin should upload his secret key to the AWS console and let S3 decrypt the objects

C S3 does not support client supplied encryption keys for server side encryption

D The admin should send the keys and encryption algorithm information with each API call

#### Explanation

Amazon S3 supports client side or server side encryption to encrypt all data at rest. The server side encryption can either have the S3 supplied AES-256 encryption key or the user can send the key along with each API call to supply his own encryption key. Amazon S3 never stores the user's encryption key. The user has to supply it for each encryption or decryption call.

<http://docs.aws.amazon.com/Ama...>

## #41

Amazon ECS

You are asked to establish a baseline for normal Amazon ECS performance in your environment by measuring performance at various times and under different load conditions. To establish a baseline, Amazon recommends that you should, at a minimum, monitor the CPU and \_\_\_ for your Amazon ECS clusters and the CPU and \_\_\_ metrics for your Amazon ECS services.

✓ memory utilization; memory reservation and utilization

B memory reservation and utilization; memory utilization

C concurrent connections; memory reservation and utilization

D memory reservation and utilization; concurrent connections

### Explanation

As you monitor Amazon ECS, store historical monitoring data so that you can compare it with current performance data, identify normal performance patterns and performance anomalies, and devise methods to address issues.

To establish a baseline, you should, at a minimum, monitor the following items:

- The CPU and memory reservation and utilization metrics for your Amazon ECS clusters
- The CPU and memory utilization metrics for your Amazon ECS services

<http://docs.aws.amazon.com/Ama...>

## #44

DynamoDB

Your system uses a multi-master, multi-region DynamoDB configuration spanning two regions to achieve high availability. For the first time since launching your system, one of the AWS Regions in which you operate went down for 3 hours, and the failover worked correctly. However, after recovery, your users are experiencing strange bugs, in which users on different sides of the globe see different data. What is a likely design issue that was not accounted for when launching?

- A** The system did not include repair logic and request replay buffering logic for post-failure, to re-synchronize data to the Region that was unavailable for a number of hours.
- B** The system did not implement DynamoDB Table Defragmentation for restoring partition performance in the Region that experienced an outage, so data is served stale.
- ✓** The system does not have Lambda Function Repair Automations, to perform table scans and check for corrupted partition blocks inside the Table in the recovered Region.
- D** The system did not use DynamoDB Consistent Read requests, so the requests in different areas are not utilizing consensus across Regions at runtime.

#### Explanation

When using multi-region DynamoDB systems, it is of paramount importance to make sure that all requests made to one Region are replicated to the other. Under normal operation, the system in question would correctly perform write replays into the other Region. If a whole Region went down, the system would be unable to perform these writes for the period of downtime. Without buffering write requests somehow, there would be no way for the system to replay dropped cross-region writes, and the requests would be serviced differently depending on the Region from which they were served after recovery.

<http://docs.aws.amazon.com/ama...>

## #47

Lambda

After having written a Lambda function which you wrote in a stateless style as required, you realise that because it is stateless it will not store some of the data that you wish to keep. Which of the following is a true statement in relation to storing persistent states when writing a Lambda function code?

- A** Any persistent state can be stored in Amazon S3, Amazon DynamoDB, or another Internet-available storage service.

B Any persistent state cannot be stored.

✓ Any persistent state can only be stored in DynamoDB.

D Any persistent state can only be stored in Amazon S3.

#### Explanation

Your Lambda function code must be written in a stateless style and must have no affinity with the underlying compute infrastructure.

Local file system access, child processes, and similar artifacts are limited to the lifetime of the request, and any persistent state should be stored in **Amazon S3, Amazon DynamoDB, or another cloud storage service.**

<https://aws.amazon.com/lambda/...>

## #51

DynamoDB

You are attempting to use the DELETE action in your DynamoDB command line using an update call. However, it is encountering an error. What might be some good things to check for before you try the action again? (Choose 3 answers)

A Check to see if the attribute you're calling DELETE is nested.

✓ Check to see if you are calling DELETE on an empty set.

✓ Check to make sure you're not calling DELETE on a number.

✓ Check the data type of the attribute you're calling DELETE on.

#### Explanation

The DELETE action only supports set data types, so you would want to check that the data type is a set. Specifying an empty set will cause an error. It also only works on top-level attributes, not nested attributes. Multiple actions can be performed in a single expression, so that would NOT cause an error.

<http://docs.aws.amazon.com/ama...>

## #52

CodeDeploy

You create a deployment configuration in AWS CodeDeploy, but you do not specify it. What happens in this scenario?

- ✓ The deployment configuration fails.
- B A popup alerts you to choose a deployment configuration.
- C AWS CodeDeploy will use a specific default deployment configuration.
- D AWS CodeDeploy will send an alarm to CloudWatch.

## Explanation

When you create a deployment in AWS CodeDeploy, you may specify an existing deployment configuration. However, if you don't specify a deployment configuration, AWS CodeDeploy will use a specific default deployment configuration.

<http://docs.aws.amazon.com/cod...>

## #53

EC2

You work for a medical research firm that is hosting their complex DNA ChIP-sequencing algorithm on one large Amazon EC2 instance. However, the algorithm takes a great deal of processing power and only a handful of users can run the algorithm on that one instance. Often, you need at least three or four instances running to accommodate for demand. However, leaving four large EC2 instances running all the time is quite expensive, especially when demand is irregular.

What might be a good solution to the problem of fluctuating demand?

- A Use AWS Lambda to write metrics out to the Lambda log. When metrics are above a certain threshold, have Amazon Simple Notification Service send out an email for the admin to add another EC2 instance.
- B Have one Auto Scaling policy (scale-out) and use AWS CloudTrail to monitor metrics, kicking-off the policy when demand is high.



- ✓ Use AWS Lambda to write metrics out to the Lambda log. When metrics are above a certain threshold, call AWS CloudTrail to trigger an Auto Scaling policy that starts up another instance.

- D Have two Auto Scaling policies (scale-in and scale-out) and use Amazon CloudWatch to kick-off the correct policy according to demand.

### Explanation

If you want to dynamically scale your instances, you need to use Auto Scaling in conjunction with Amazon CloudWatch. In this case, you want to scale out when you have high demand, and scale in when demand drops, so you would need two Auto Scaling policies that fire based on metrics from CloudWatch.

<http://docs.aws.amazon.com/aut...>

## #54

SQS

A user has created a queue named "myqueue" in US-East region with AWS SQS. The user's AWS account ID is 123456789012. If the user wants to perform some action on this queue, which of the below Queue URLs should he use?

- A <http://sqs.123456789012.us-east-1.amazonaws.com/myqueue>

- B <http://123456789012.sqs.us-east-1.amazonaws.com/myqueue>

- C <http://sqs.us-east-1.amazonaws.com/123456789012/myqueue>

- ✓ <http://sqs.amazonaws.com/123456789012/myqueue>

### Explanation

When creating a new queue in SQS, the user must provide a queue name that is unique within the scope of all queues of user's account. If the user creates queues using both the latest WSDL and a previous version, he will have a single namespace for all his queues. Amazon SQS assigns each queue created by user an identifier called a queue URL, which includes the queue name and other components that Amazon SQS determines. Whenever the user wants to perform an action on a queue, he must provide its queue URL. The queue URL for the account id 123456789012 & queue name "myqueue" in US-East-1 region will be <http://sqs.us-east-1.amazonaws.com/123456789012/myqueue>.

<http://docs.aws.amazon.com/AWS...>

#55

Lambda

You are writing a Lambda function for a mobile app and want to control how it is invoked (synchronous or asynchronous). What are your options for controlling the type of invocation? (Choose 2 answers)

✓ You can manually invoke the function via the AWS Command-line.

✓ You can change the settings for Amazon Cognito and specify which invocation type you would like.

C You can write code in you mobile app that calls the Lambda function directly.

D You can manually adjust your Amazon DynamoDB stream to reflect your preferred invocation type.

#### Explanation

When you are using AWS services as event sources, the invocation type is predetermined for each of these services. Therefore, you can only control the invocation type when you do it manually from the command-line or directly from code.

<http://docs.aws.amazon.com/lam...>

#58

EC2

You have recently configured an AWS VPC for a new client. The VPC consists of several EC2 instances in a public subnet. An IT administrator for one of your clients has called you about an inability to connect to an EC2 instance. He is trying to connect from a Windows laptop using Putty for Windows and the .pem file you provided him along with instructions for using PuTTY and PuTTYgen. What could be the problem?

A The RDP port is not open.

B The .pem file needs converted to a .ppk file.

✓ There is no Virtual Private Gateway.

D The admin is not configured in IAM.

#### Explanation

If you use PuTTY to connect to your instances verify that your private key (.pem) file has been converted to the format recognized by PuTTY (.ppk). PuTTY does not natively support the private key format (.pem) generated by Amazon EC2. PuTTY has a tool named PuTTYgen, which can convert keys to the required PuTTY format (.ppk). You must convert your private key into this format (.ppk) before attempting to connect to your instance using PuTTY.

<http://docs.aws.amazon.com/AWS...>

## #59

KMS

When Amazon S3 uses AWS KMS, it uses the Amazon S3 encryption client in the AWS SDK from your own application to encrypt objects and upload them to Amazon S3. What does this method allow you to do?

A It allows you to encrypt your data whenever you want.

✓ It allows you to encrypt your data with low latency.

C It allows you to decrypt your data using the Amazon S3 service.

D It allows you to encrypt your data locally to ensure its security as it passes to the Amazon S3 service.

#### Explanation

When Amazon S3 uses AWS KMS, it uses the Amazon S3 encryption client in the AWS SDK from your own application to encrypt objects and upload them to Amazon S3. This method allows you to encrypt your data locally to ensure its security as it passes to the Amazon S3 service. The S3 service receives your encrypted data and does not play a role in encrypting or decrypting it.

<http://docs.aws.amazon.com/kms...>

#61

AWS CodeCommit

How are repositories in AWS CodeCommit encrypted?

- A Repositories are encrypted with the RSA public-key encryption algorithm.
- ✓ Repositories are encrypted by enabling the AWS Key Management Service (KMS).
- C Repositories are automatically encrypted at rest.
- D Repositories are encrypted by providing IAM access to the AWS Key Management Service (KMS).

**Explanation**

In AWS CodeCommit, repositories are automatically encrypted at rest. No customer action is required. AWS CodeCommit uses AWS Key Management Service (KMS) to perform this encryption.

<http://aws.amazon.com/codecomm...>

#62

ElastiCache

What is one reason that AWS does not recommend that you configure your ElastiCache so that it can be accessed from outside AWS?

- A The NAT instance serves as a single point of failure.
- B The performance of the ElastiCache cluster is no longer controllable.
- ✓ The metrics reported by CloudWatch are more difficult to report.
- D The ElastiCache cluster becomes more prone to failures.

**Explanation**

One of the reasons that AWS does not recommend that you configure your ElastiCache so that it can be accessed from outside AWS is that the NAT instance serves as a single point of failure. Here are other reasons that this practice is not recommended:

- The NAT instance is acting as a proxy between clients and multiple clusters. The addition of a proxy impacts the performance of the cache cluster. The impact increases with number of cache clusters you are accessing through the NAT instance.
- The traffic from clients to the NAT instance is unencrypted. Therefore, you should avoid sending sensitive data via the NAT instance.
- The NAT instance adds the overhead of maintaining another instance.

<http://docs.aws.amazon.com/Ama...>

## #63

EC2   OpsWorks

Your customer is currently manually deploying EC2 instances to a VPC and provisioning them to serve various roles in their web application. The customer has a limited operations staff and developers are spending a significant amount of time installing packages & frameworks and configuring software. The customer would like to automate the installation of packages & frameworks and software configuration as well as scale EC2 deployments based on traffic to allow the developers to focus on development as opposed to infrastructure. There is a preference to minimize cost for any solution implemented. Which solution should you recommend?

**A**   OpsWorks Stacks

**B**   OpsWorks for Chef Automate

✓ **Kinesis**

**D**   CloudWatch

**Explanation**

OpsWorks stacks allows for automation of package installation, installation of programming language and frameworks, and software configuration. It also allows you to scale based on changing traffic levels and is offered with no additional charge for use with EC2 (you only pay for the resources created). OpsWorks for Chef Automate has costs associated with the number of nodes connected to the Chef server and the underlying instance running the Chef Server. CloudWatch is a monitoring solution. Kinesis is used to collect and load streaming data.

<https://aws.amazon.com/opswork...>

## #64

Elastic Beanstalk

Which of the following statements holds true during rolling updates in Elastic Beanstalk?

- A Elastic Beanstalk applies a separate timeout to each batch in the operation as a default, but this can be changed.
- B Elastic Beanstalk applies a separate timeout to each batch in the operation.
- C Elastic Beanstalk ignores timeouts.
- ✓ Elastic Beanstalk applies the same timeout to each batch in the operation.

**Explanation**

During rolling updates, Elastic Beanstalk will apply a separate timeout to each batch in the operation. The timeout is also set as part of the environment's rolling update configuration. If all instances in the batch are healthy within the command timeout, the operation will continue to the next batch. If not, the operation fails.

<http://docs.aws.amazon.com/ela...>

## #65

EC2

You are a DevOps engineer responsible for supporting your company's AWS infrastructure, consisting of multiple EC2 instances running in a VPC, DynamoDB, SQS, and S3. You are working on provisioning a new S3 bucket, which will ultimately contain sensitive data. How can you encrypt that data in-flight, into, and out of S3? (Choose 2 answers)

- A Set the server-side encryption option after upload.
- ✓ Use the encrypted SSL/TLS endpoint.

✓ Set the server-side encryption option on upload.

D Encrypt it on the client-side before uploading.

Explanation

To encrypt your S3 objects in-flight, you need to use the TLS endpoint; alternatively, you can encrypt the data yourself on the client side before upload.

<http://docs.aws.amazon.com/Ama...>

Learning Path

SOLUTIONS

- Manage the Multi-Cloud
- Migrate to the Cloud
- Develop Tech Talent
- Optimize Cloud Spend

TRAINING LIBRARY

- Amazon Web Services
- Microsoft Azure
- Google Cloud Platform
- DevOps
- Cloud Migration
- Security
- Serverless
- Big Data
- Machine Learning
- Cloud Computing Fundamentals
- Certifications
- Partner Ecosystem
- Upcoming Training

RESOURCES

- Blog
- Events
- Newsroom
- Webinars
- Case Studies
- Whitepapers
- Support
- AWS Certifications
- Azure Certifications
- Google Cloud Certifications
- Courses
- Release Notes

COMPANY

- About
- Careers We are hiring!
- Partners
- Contact Us
- Instructors



Copyright © 2019 Cloud Academy Inc. All rights reserved.

[Terms of Use](#) [Privacy Policy](#) [Sitemap](#) [System Status](#)

